

# Training ML models

Cambridge ICCS summer school  
[cambridge-iccs.github.io/summerschool](https://cambridge-iccs.github.io/summerschool)

Will Handley

2022-09-22

Slides & installation instructions at  
<https://github.com/handley-lab/2022-cambridge-iccs>

# Overview

- ▶ Title is rather broad, and we have  $< 2h$ , the morning after the conference dinner.
- ▶ Aim to build a framework of **understanding** in the context of a few examples
- ▶ Should be able to answer:
  - ▶ What is an ML model?
  - ▶ How do I avoid pitfalls in training them?
  - ▶ Which resources should I reach for in the future?

# The Machine Learning Python stack

## 0. numpy

- ▶ Layer zero – vector maths & array-based programming
- ▶ Advanced users: broadcasting,  $x @ y$ ,  $z[:, \text{None}]$ , ufuncs

## 1. scipy & pandas

- ▶ Extends numpy to numerical algorithms and excel-like array functionality
- ▶ pandas often a data scientist's weakest point


## 2. scikit-learn

- ▶ Entry-level machine learning
- ▶ Extends to allow estimators, transformers & predictors
- ▶ With a few key concepts this is a consistent and versatile ML framework

## 3. Keras/TensorFlow PyTorch,

- ▶ Deep learning tools
- ▶ Familiarity with the previous layers greatly enhances effectiveness


+ matplotlib for plotting (others exist, extending and reducing flexibility)



# What is machine learning?

# What is machine learning?

- ▶ A computer program which can program itself to perform a task
  - ▶ Problems with lots of tuning/rules
  - ▶ Problems with no traditional solution
  - ▶ Fluctuating environments
  - ▶ Gaining insight about complex data
- ▶ Traditional programs
  - ▶ Quicksort
  - ▶ Pong
- ▶ Machine learning
  - ▶ Spam filter
  - ▶ Netflix suggestions
  - ▶ Speech recognition
  - ▶ Dall-E



# The machine learning pipeline

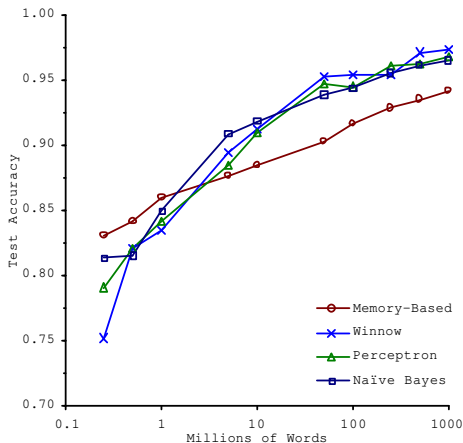
1. Problem framing
2. Data acquisition
3. Visualisation
4. Data preparation/munging
5. Selecting & training a model
6. Tuning a model
7. Launch, monitor & maintain

# The machine learning pipeline

1. Problem framing
  - ▶ Big picture
  - ▶ Selecting performance measures/objectives
  - ▶ Checking assumptions/bias
2. Data acquisition
3. Visualisation
4. Data preparation/munging
5. Selecting & training a model
6. Tuning a model
7. Launch, monitor & maintain

# The machine learning pipeline

1. Problem framing
2. Data acquisition
  - ▶ Gather your data
  - ▶ Selecting performance measures/objectives
  - ▶ Checking assumptions/bias
3. Visualisation
4. Data preparation/munging
5. Selecting & training a model
6. Tuning a model
7. Launch, monitor & maintain



[doi:10.3115/1073012.1073017]



# The machine learning pipeline

1. Problem framing
2. Data acquisition
3. Visualisation
  - ▶ using pandas+matplotlib skills to explore idiosyncrasies of the data
4. Data preparation/munging
5. Selecting & training a model
6. Tuning a model
7. Launch, monitor & maintain

# The machine learning pipeline

1. Problem framing
2. Data acquisition
3. Visualisation
4. Data preparation/munging
  - ▶ using scikit-learn to clean & transform the data
5. Selecting & training a model
6. Tuning a model
7. Launch, monitor & maintain

# The machine learning pipeline

1. Problem framing
2. Data acquisition
3. Visualisation
4. Data preparation/munging
5. Selecting & training a model
  - ▶ Topic of this session
6. Tuning a model
7. Launch, monitor & maintain

# The machine learning pipeline

1. Problem framing
2. Data acquisition
3. Visualisation
4. Data preparation/munging
5. Selecting & training a model
6. Tuning a model
  - ▶ Topic of this session
7. Launch, monitor & maintain

# The machine learning pipeline

1. Problem framing
2. Data acquisition
3. Visualisation
4. Data preparation/munging
5. Selecting & training a model
6. Tuning a model
7. Launch, monitor & maintain
  - ▶ For researchers this could be github distributing
  - ▶ For industry this would mean real-world shipping

# Categories of machine learning

## Supervised

Regression  
Classification

## Unsupervised

Clustering, Visualisation,  
Dimensionality reduction

## Semisupervised

Google Photos

## Reinforcement

AlphaGo, GANNS

# Categories of machine learning

## Supervised

Regression  
Classification

## Unsupervised

Clustering, Visualisation,  
Dimensionality reduction

## Semisupervised

Google Photos

## Reinforcement

AlphaGo, GANNS

## Batch

Offline learning using all available data.

## Online

Training/updating on-the-fly on  
mini-batches, for  
memory-bound/out-of-core

# Categories of machine learning

## Supervised

Regression  
Classification

## Unsupervised

Clustering, Visualisation,  
Dimensionality reduction

## Semisupervised

Google Photos

## Reinforcement

AlphaGo, GANNS

## Batch

Offline learning using all available data.

## Online

Training/updating on-the-fly on  
mini-batches, for  
memory-bound/out-of-core

## Instance based learning

“Learn-by-heart” – given a similarity  
measure, compare/regress/classify new  
examples onto existing data

## Model-based learning

Build a parameterised model of data, **train**  
it, then make predictions



# Challenges of Machine Learning

- ▶ Bad data
  - ▶ Not enough
  - ▶ Not representative
  - ▶ Poor quality (outliers, noise)
  - ▶ Irrelevant/Poor features
- ▶ Bad algorithm/Bad training (Focus of this workshop)
  - ▶ Overfitting
  - ▶ Underfitting

## Example 1: Introduction to regression

Slides & installation instructions at  
<https://github.com/handley-lab/2022-cambridge-iccs>

# Training an ML model

- ▶ Three ingredients to training

1. Input data/features  $\mathbf{x}^{(i)}$ , output data  $y^{(i)}$ , where  $i = 1, \dots, n_{\text{obs}}$
2. Parameterised model  $y = h_{\theta}(\mathbf{x})$ , where  $h$  is the model and  $\theta$  are its parameters
3. Loss function(s)  $L(y_{\text{pred}}, y)$

- ▶ Train the parameters on a training subset by solving the mathematical problem

$$\hat{\theta} = \min_{\theta} \sum_{i \in \text{train}} L_{\text{train}}(h_{\theta}(\mathbf{x}^{(i)}), y^{(i)})$$

- ▶ Choose the best model by minimising a (possibly different) loss on a validation subset

$$\hat{h} = \min_h \sum_{i \in \text{validation}} L_{\text{validation}}(h_{\hat{\theta}}(\mathbf{x}^{(i)}), y^{(i)})$$

- ▶ Finally test the best model on set-aside testing data.

# 1. Data

- ▶ The initial data  $\{(\mathbf{x}^{(i)}, y^{(i)}), i = 1, \dots, n_{\text{obs}}\}$  must usually be transformed
- ▶ Relevant features should be selected
- ▶ Relevant combinations of features should be considered (e.g. computing rates/sums)
- ▶ This is known as feature engineering
- ▶ Features should then be normalised, Either:
  - `min-max scaled` data lie in  $[0, 1]$  (`sklearn.preprocessing.MinMaxScaler`)
  - `standardised` data have mean 0 and std 1 (`sklearn.preprocessing.StandardScaler`)
- ▶ The `sklearn` way to do this is to chain a set of these transformations together in a `sklearn.pipeline.Pipeline`
- ▶ This is essential since almost all machine learning algorithms are not covariant, and will fail on unnormalised data.

## 2. Models

- ▶ “How to choose models” would fill a whole other session.
- ▶ You may recognise some of the standard choices.
- ▶ Models have trainable parameters  $\theta$ , and hyperparameters.

### Supervised

- ▶ *k*-Nearest Neighbors
- ▶ Linear Regression
- ▶ Logistic Regression
- ▶ Support Vector Machines
- ▶ Decision Trees and Random Forests
- ▶ Neural Networks

### Semisupervised

- ▶ deep belief networks
- ▶ RBMs

### Reinforcement

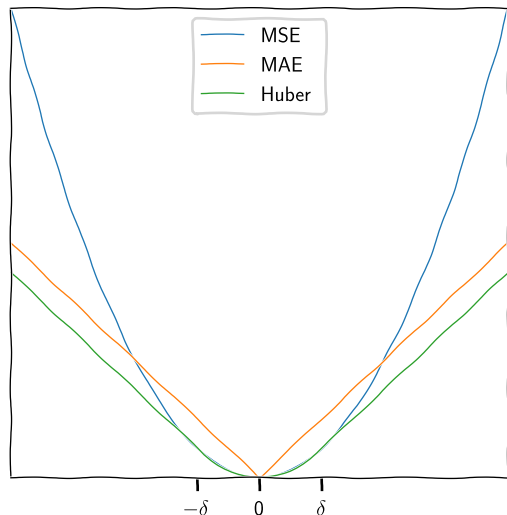
- ▶ AlphaGo
- ▶ GANs

### Unsupervised

- ▶ Clustering
  - ▶ *K*-Means
  - ▶ DBSCAN
  - ▶ Hierarchical Cluster Analysis
- ▶ Anomaly detection
  - ▶ One-class SVM
  - ▶ Isolation forest
- ▶ Visualisation and dimensionality reduction
  - ▶ (Kernel) PCA
  - ▶ Locally-Linear Embedding
  - ▶ t-SNE
- ▶ Association rule learning
  - ▶ Apriori & Eclat

### 3. Loss functions

- ▶ All that is needed is something which measures how “close” a model’s prediction is to the true answer
- ▶ The loss function you train on does not need to be the same as the testing/validation metric.
- ▶ Mean square error (MSE)  $L = (\Delta y)^2$ 
  - ▶ smooth (differentiable)
- ▶ Mean absolute error (MAE)  $L = |\Delta y|$ 
  - ▶ robust to outliers
- ▶ Huber loss
$$L = \begin{cases} \frac{1}{2}(\Delta y)^2 & |\Delta y| < \delta \\ \delta(|\Delta y| - \frac{1}{2}\delta) & |\Delta y| \geq \delta \end{cases}$$
  - ▶ combines benefits of both



# Principles of Training, Validation & Testing splits

- ▶ The sure-fire way to know how a model will generalise is to hold back data for testing.
- ▶ We therefore split data into three categories

## Training

~80% of the data.  
Used for learning parameters.

## +Validation

Used for learning  
hyperparameters.

## Testing

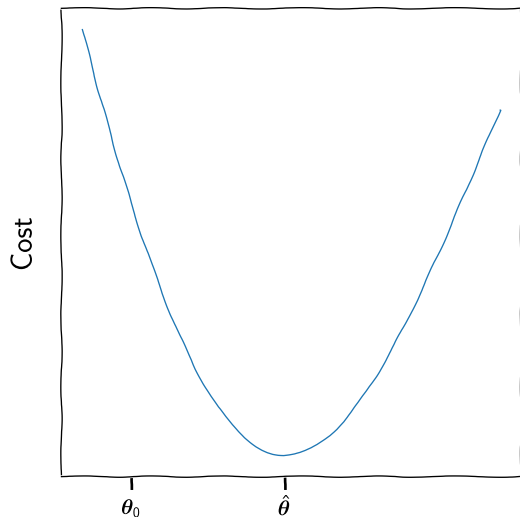
Holdout ~20% of the data.  
Should be ideally “one-shot”.

- ▶ “Validation set” is also known as the “development/dev set”.
- ▶ If we do repeated holdout validation on many small validation sets, this is **cross-validation**.
- ▶ Gotchas:
  - ▶ standardisation/data preparation should only use the training+validation set
    - ▶ Failing to do this reduces generalisability
  - ▶ ideally the splitting procedure should be random, but seedable
    - ▶ Failing to do this reduces repeatability & reliability
  - ▶ `sklearn.model_selection.StratifiedKFold` accomplishes these and more

# Gradient descent

- To minimise a function  $f(\theta)$ , start somewhere  $\theta_0$  and go downhill (down the gradient) by some step  $\eta$

$$\theta_{k+1} = \theta_k - \eta \nabla_{\theta} f(\theta)$$

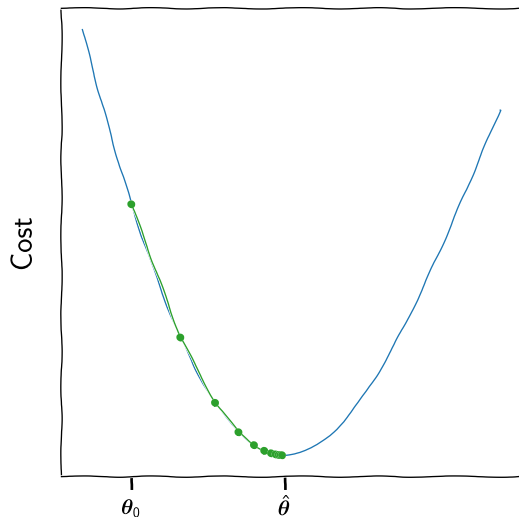




# Gradient descent

- To minimise a function  $f(\theta)$ , start somewhere  $\theta_0$  and go downhill (down the gradient) by some step  $\eta$

$$\theta_{k+1} = \theta_k - \eta \nabla_{\theta} f(\theta)$$

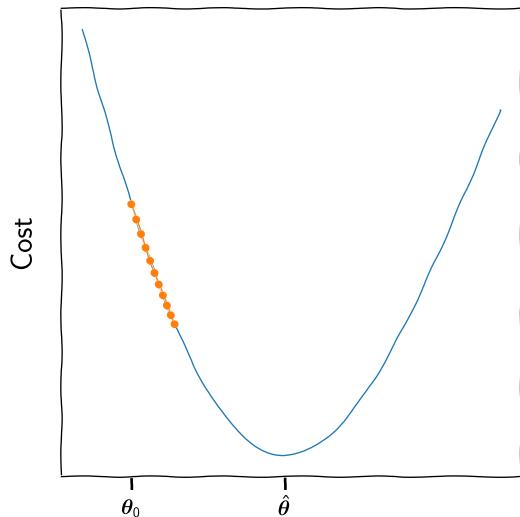


# Gradient descent

- ▶ To minimise a function  $f(\theta)$ , start somewhere  $\theta_0$  and go downhill (down the gradient) by some step  $\eta$

$$\theta_{k+1} = \theta_k - \eta \nabla_{\theta} f(\theta)$$

- ▶ Problems:
  - ▶ Learning rate too slow



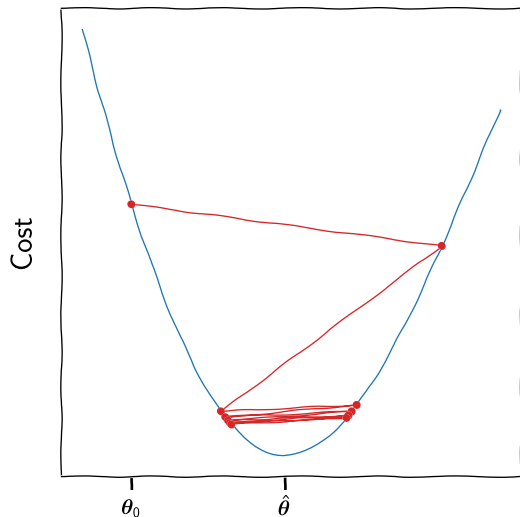
# Gradient descent

- ▶ To minimise a function  $f(\theta)$ , start somewhere  $\theta_0$  and go downhill (down the gradient) by some step  $\eta$

$$\theta_{k+1} = \theta_k - \eta \nabla_{\theta} f(\theta)$$

- ▶ Problems:

- ▶ Learning rate too slow
- ▶ Learning rate too fast




# Gradient descent

- ▶ To minimise a function  $f(\theta)$ , start somewhere  $\theta_0$  and go downhill (down the gradient) by some step  $\eta$

$$\theta_{k+1} = \theta_k - \eta \nabla_{\theta} f(\theta)$$

- ▶ Problems:

- ▶ Learning rate too slow
- ▶ Learning rate too fast
- ▶ Local minima



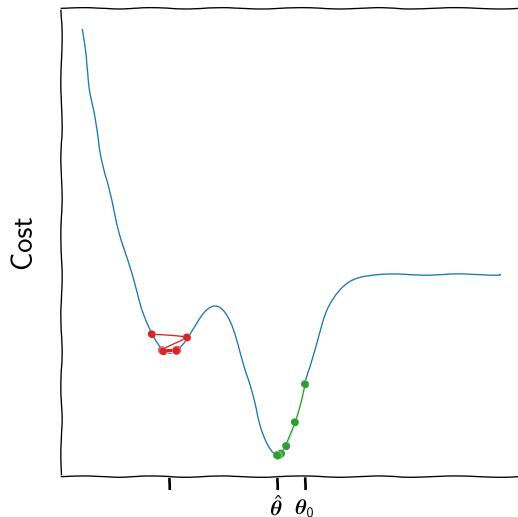
# Gradient descent

- ▶ To minimise a function  $f(\theta)$ , start somewhere  $\theta_0$  and go downhill (down the gradient) by some step  $\eta$

$$\theta_{k+1} = \theta_k - \eta \nabla_{\theta} f(\theta)$$

- ▶ Problems:

- ▶ Learning rate too slow
- ▶ Learning rate too fast
- ▶ Local minima



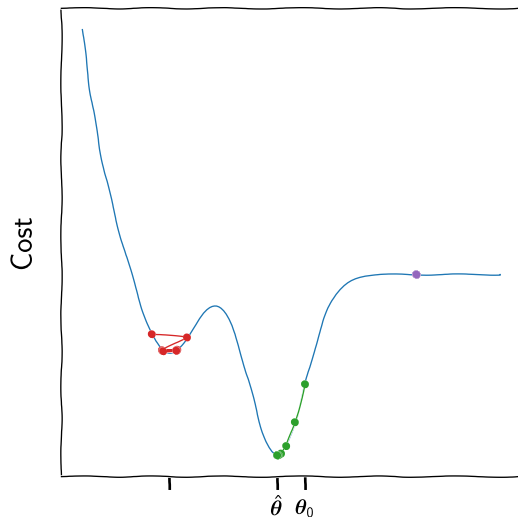
# Gradient descent

- ▶ To minimise a function  $f(\theta)$ , start somewhere  $\theta_0$  and go downhill (down the gradient) by some step  $\eta$

$$\theta_{k+1} = \theta_k - \eta \nabla_{\theta} f(\theta)$$

- ▶ Problems:

- ▶ Learning rate too slow
- ▶ Learning rate too fast
- ▶ Local minima
- ▶ stalling



# Gradient descent

- ▶ To minimise a function  $f(\theta)$ , start somewhere  $\theta_0$  and go downhill (down the gradient) by some step  $\eta$

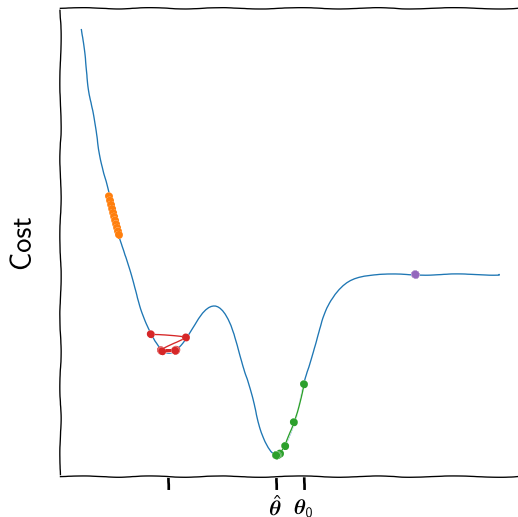
$$\theta_{k+1} = \theta_k - \eta \nabla_{\theta} f(\theta)$$

- ▶ Problems:

- ▶ Learning rate too slow
- ▶ Learning rate too fast
- ▶ Local minima
- ▶ stalling

- ▶ Choosing the learning rate, or more generally tuning the learning schedule can be the hardest part of training.

- ▶ All much harder in higher dimensions



## Example 2: Machine learning with scikit-learn

Go back to the notebook `training_ml_models.ipynb`

For more detail on the scikit-learn API (beyond the docs) I recommend [arxiv:1309.0238]



# How we get the gradients: autodiff

- ▶ There are three ways to get a computer to compute a gradient  $\nabla_{\theta} f(\theta)$ 
  1. Analytically
    - ▶ Painstakingly coding the function explicitly
    - ▶ Accurate, but practically impossible for all but the simplest functions
  2. Numerically
    - ▶ Computing finite differences  $[\nabla_{\theta} f(\theta)]_i \approx [f(\theta + \delta \hat{e}_i) - f(\theta)] / \delta$
    - ▶ Easy, but prone to numerical instability
    - ▶ Expensive – costs  $\sim \mathcal{O}(n)$  function evaluation for each coordinate direction  $\hat{e}_i$
  3. Automatically
    - ▶ Every computer programme is composed of logic, multiplication and addition.
    - ▶ A smart enough computer, equipped with a chain rule, can therefore differentiate any code.
    - ▶ Remarkably, this can be done at the same cost as computing  $f$
- ▶ Autodiff has been around for a long time, but unless the programming language is designed with it in mind it is difficult in practice.
- ▶ Modern ML codes are.
- ▶ In traditional ML literature this is wrapped up in a mythology of “backpropagation” – equivalent to “reverse mode autodiff”, well suited to high-chaining with few outputs.

# Variations on gradient descent

## Batch gradient descent (full)

Compute the gradient in full at each step:

$$\nabla_{\theta} \sum_i \left[ L(h_{\theta}(\mathbf{x}^{(i)}), y^{(i)}) \right]$$

Deterministic.

Slow for large data sets.

## Stochastic gradient descent (SGD)

Compute the gradient on one random data point

$$\nabla_{\theta} \left[ L(h_{\theta}(\mathbf{x}^{(j)}), y^{(j)}) \right]$$

Blazingly fast.

Randomness can escape minima.

## Mini-batch gradient descent

Compute the gradient on a (small) subset of points

$$\nabla_{\theta} \sum_{i \in \text{batch}_j} \left[ L(h_{\theta}(\mathbf{x}^{(i)}), y^{(i)}) \right]$$

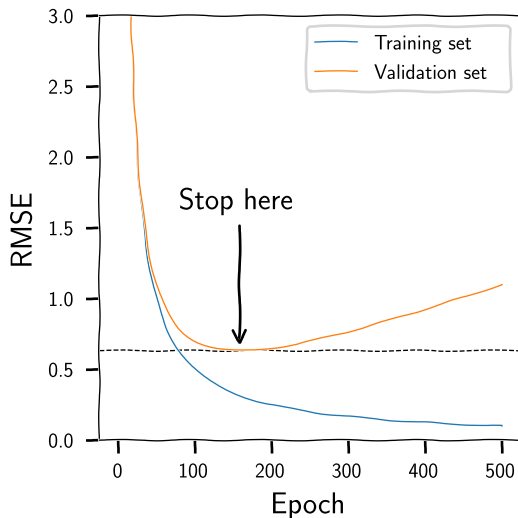
GPU accelerable.

Batch size  $< 32$ .

- In addition SGD & mini-batch have benefits for out-of-core learning, when the data are too large to fit into memory.

# Regularisation

- ▶ Can reduce overfitting by constraining the degrees of freedom of the model
  - ▶ reduce the number of free parameters
  - ▶ constrain the parameters to a reduced range
- ▶ Regularisation (only applied at training)
  - ▶ Ridge regression:  $L(\theta) + \frac{1}{2}\alpha \sum_i \theta_i^2$
  - ▶ Lasso regression:  $L(\theta) + \alpha \sum_i |\theta_i|$ 
    - ▶ Least absolute shrinkage & selection operator
    - ▶ Promotes sparsity
  - ▶ Elastic Net:  $L(\theta) + r\alpha \sum_i |\theta_i| + \frac{1-r}{2}\alpha \sum_i \theta_i^2$
- ▶ Early stopping
  - ▶ Halt training when validation error increases.
  - ▶ As close to a free lunch as one gets.
  - ▶ Very popular.



# Neural networks

A multi-layer perceptron (MLP)  
is mathematically

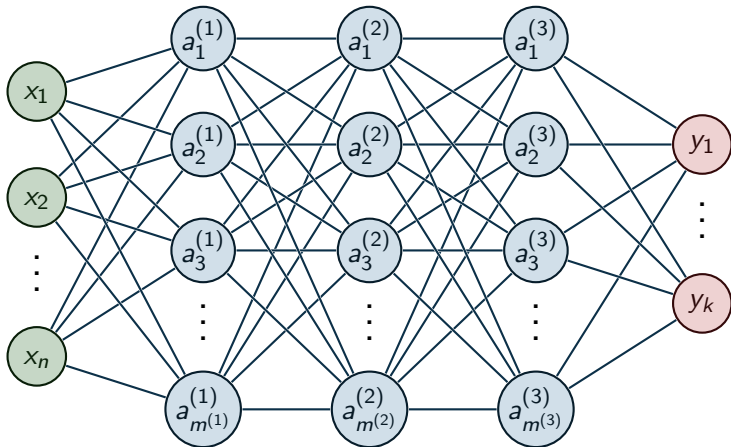
$$a_i^{(\ell+1)} = \phi_i^{(\ell)} \left( \sum_{j=1}^{m^{(\ell)}} w_{ij}^{(\ell)} a_j^{(\ell)} + b_i^{(\ell)} \right)$$

or written in vectors:

$$\mathbf{a}^{(\ell+1)} = \Phi^{(\ell)} \left( \mathbf{W}^{(\ell)} \mathbf{a}^{(\ell)} + \mathbf{b}^{(\ell)} \right)$$

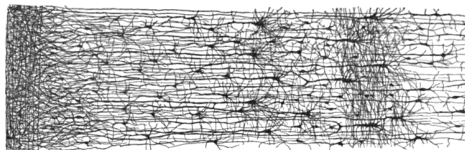
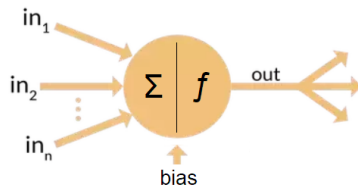
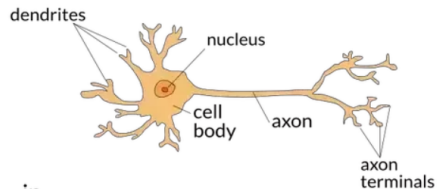
$$\mathbf{a}^{(0)} = \mathbf{x}, \quad \mathbf{y} = \mathbf{a}^{(m+1)}$$

or graphically:



# Neural networks & “deep” learning

- ▶ NNs were originally inspired by biology.
- ▶ Old technology (1960s), came of age in 1990s.
  - ▶ rise in computing power
  - ▶ innovations in training
  - ▶ funding/interest from tech companies
- ▶ Universal approximation theorem:  
Any function  $\mathbb{R}^n \rightarrow \mathbb{R}^k$  can be approximated by a sufficiently wide single hidden layer NN
- ▶ So why do we need “deep” (multi-layer) networks?
- ▶ Earlier layers perform **feature learning** to pipe into final universal approximating layer
- ▶ Enables the rudiments of **transfer learning**



# Deep learning tools

## TensorFlow

(Google/Alphabet)

- ▶ 2015
- ▶ Symbolic math library
- ▶ Keras makes easier
- ▶ More popular in industry



## PyTorch

(Facebook/Meta)

- ▶ 2017
- ▶ Easier to get started
- ▶ Faster than Keras
- ▶ More popular in research



- ▶ Keras is a Python API to TensorFlow, CNTK & Theano
- ▶ CNTK is MicroSoft's (now defunct) "cognitive toolkit"
- ▶ Theano used to be a giant in the field, and is a 2007 "grandfather" to the rest. Now only used by research/legacy code

**In summary:** the big boys & girls in industry use TensorFlow, but since PyTorch is preferred and developed by research now, it may become dominant in a few years time. Keras is a wrapper to TensorFlow which makes it more native, but slower and less flexible

# Neural Network Anatomy & Training

The dials you can twiddle

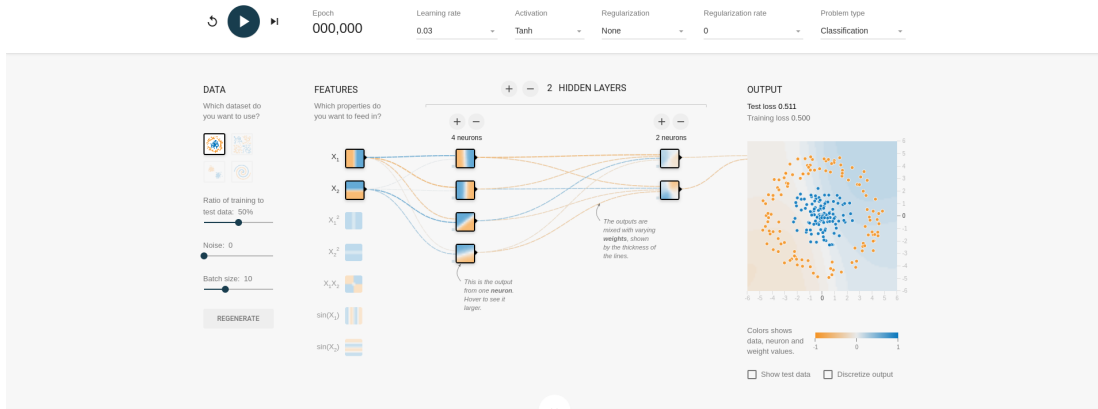
## Anatomy

- ▶ Number of hidden layers
- ▶ Width of hidden layers
- ▶ Activation functions

## Training

- ▶ Loss function
- ▶ Optimiser
- ▶ Initialisation
- ▶ Normalisation
- ▶ Regularisation
- ▶ Learning rate schedule

Tinker With a **Neural Network** Right Here in Your Browser.  
Don't Worry, You Can't Break It. We Promise.



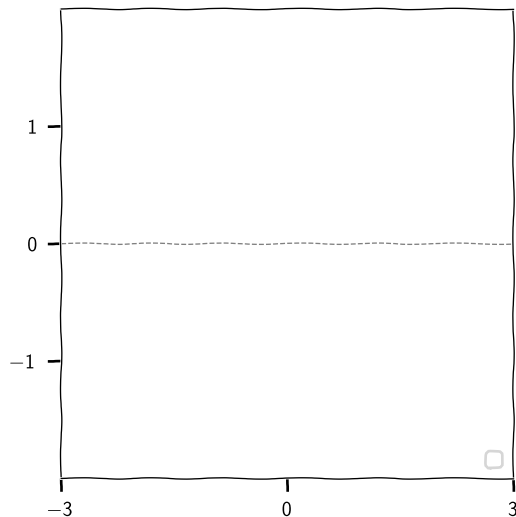


# Hyperparameter tuning

- ▶ The amount of things you can change in a neural network is both an advantage and the main drawback
- ▶ Coming up with new combinations forms a thriving field in itself
- ▶ There are some automated tools for hyperparameter tuning beyond the cross-validation grids we've found before
  - ▶ Hyperopt (Python)
  - ▶ Hyperkeras, kopt or talos (Keras)
  - ▶ Scikit-Optimise (e.g. BayesSearchCV)
  - ▶ Spearmint (more Bayesian optimisation)
  - ▶ Sklearn-Deap (evolutionary algorithms)
  - ▶ ...
- ▶ There are also cloud computing services (Google Cloud, Arimo, SigOpt, Oscar, ...)
- ▶ Finding a more principled way to choose these parameters is a big unsolved problem.
- ▶ Nevertheless, there are still some principles which are helpful to know

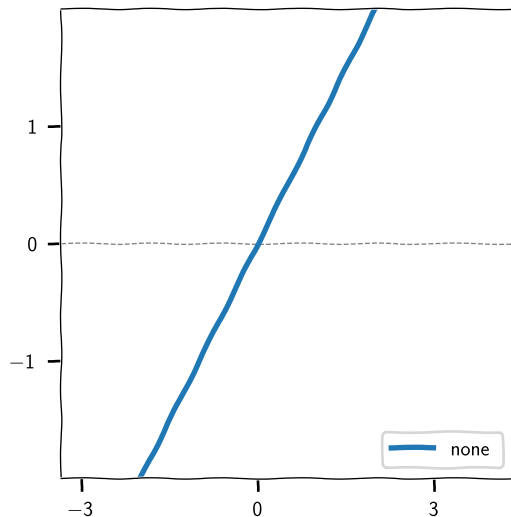
# Activation functions

- ▶ There is now a veritable zoo of activation functions
- ▶ Important properties to consider:
  - ▶ Smoothness
  - ▶ Saturation (at either end)
- ▶ Different roles depending on layer
  - ▶ At output it is useful to e.g. impose positivity with softmax or  $[0,1]$  boundedness with logistic
  - ▶ For inner layers one may want symmetry/infinite range



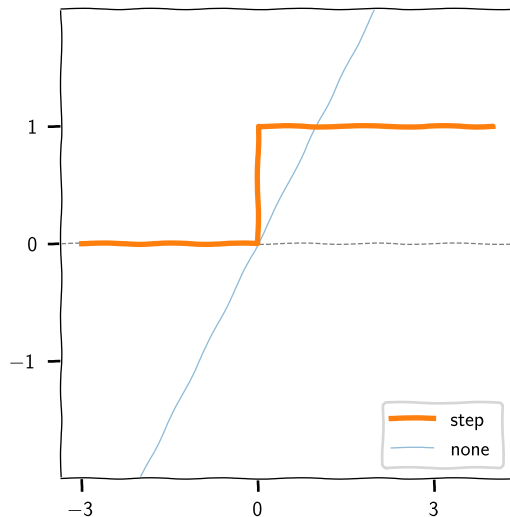
# Activation functions

- ▶ There is now a veritable zoo of activation functions
- ▶ Important properties to consider:
  - ▶ Smoothness
  - ▶ Saturation (at either end)
- ▶ Different roles depending on layer
  - ▶ At output it is useful to e.g. impose positivity with softmax or  $[0,1]$  boundedness with logistic
  - ▶ For inner layers one may want symmetry/infinite range



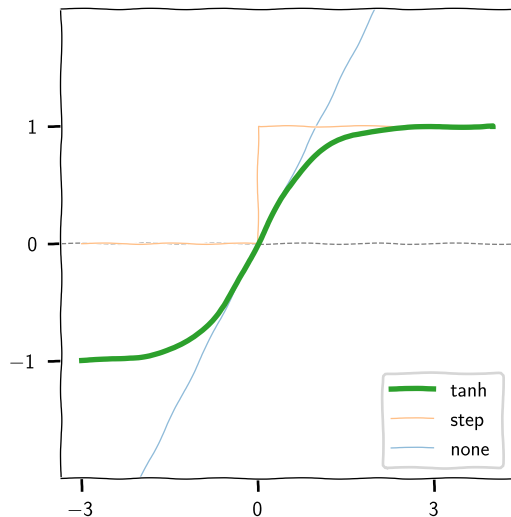
# Activation functions

- ▶ There is now a veritable zoo of activation functions
- ▶ Important properties to consider:
  - ▶ Smoothness
  - ▶ Saturation (at either end)
- ▶ Different roles depending on layer
  - ▶ At output it is useful to e.g. impose positivity with softmax or  $[0,1]$  boundedness with logistic
  - ▶ For inner layers one may want symmetry/infinite range



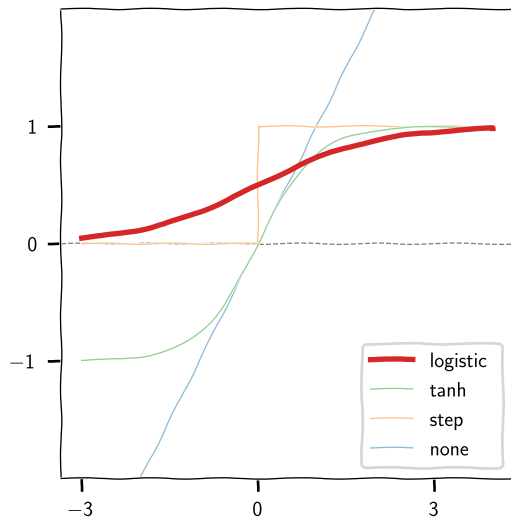
# Activation functions

- ▶ There is now a veritable zoo of activation functions
- ▶ Important properties to consider:
  - ▶ Smoothness
  - ▶ Saturation (at either end)
- ▶ Different roles depending on layer
  - ▶ At output it is useful to e.g. impose positivity with softmax or  $[0,1]$  boundedness with logistic
  - ▶ For inner layers one may want symmetry/infinite range



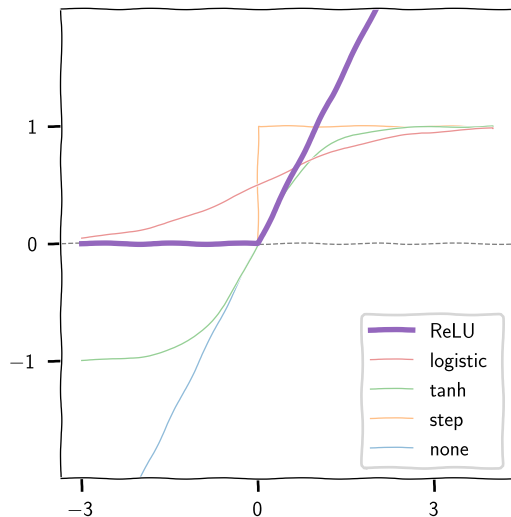
# Activation functions

- ▶ There is now a veritable zoo of activation functions
- ▶ Important properties to consider:
  - ▶ Smoothness
  - ▶ Saturation (at either end)
- ▶ Different roles depending on layer
  - ▶ At output it is useful to e.g. impose positivity with softmax or  $[0,1]$  boundedness with logistic
  - ▶ For inner layers one may want symmetry/infinite range



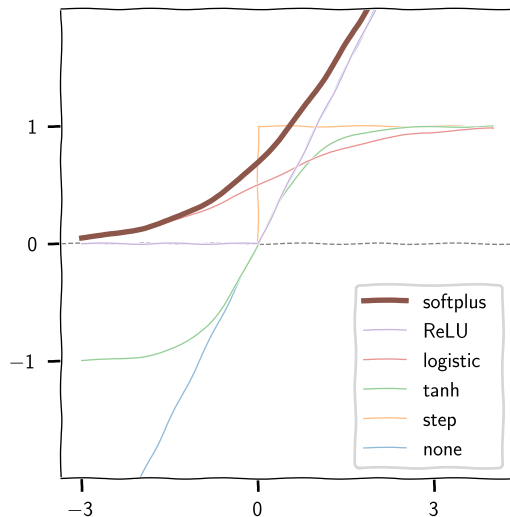
# Activation functions

- ▶ There is now a veritable zoo of activation functions
- ▶ Important properties to consider:
  - ▶ Smoothness
  - ▶ Saturation (at either end)
- ▶ Different roles depending on layer
  - ▶ At output it is useful to e.g. impose positivity with softmax or  $[0,1]$  boundedness with logistic
  - ▶ For inner layers one may want symmetry/infinite range



# Activation functions

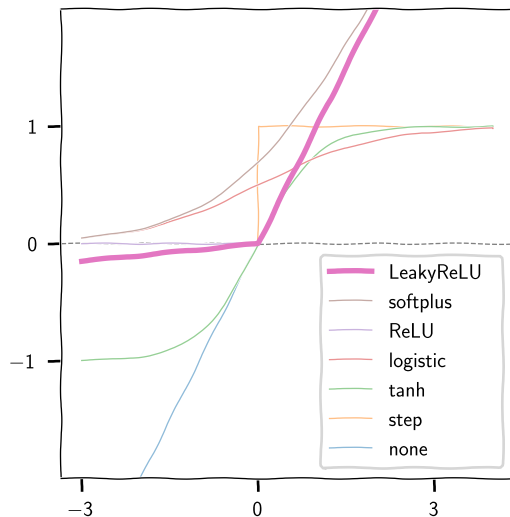
- ▶ There is now a veritable zoo of activation functions
- ▶ Important properties to consider:
  - ▶ Smoothness
  - ▶ Saturation (at either end)
- ▶ Different roles depending on layer
  - ▶ At output it is useful to e.g. impose positivity with softmax or  $[0,1]$  boundedness with logistic
  - ▶ For inner layers one may want symmetry/infinite range





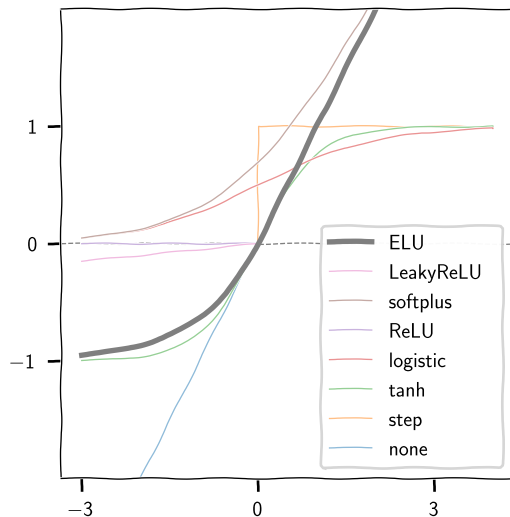
# Activation functions

- ▶ There is now a veritable zoo of activation functions
- ▶ Important properties to consider:
  - ▶ Smoothness
  - ▶ Saturation (at either end)
- ▶ Different roles depending on layer
  - ▶ At output it is useful to e.g. impose positivity with softmax or  $[0,1]$  boundedness with logistic
  - ▶ For inner layers one may want symmetry/infinite range




# Activation functions

- ▶ There is now a veritable zoo of activation functions
- ▶ Important properties to consider:
  - ▶ Smoothness
  - ▶ Saturation (at either end)
- ▶ Different roles depending on layer
  - ▶ At output it is useful to e.g. impose positivity with softmax or  $[0,1]$  boundedness with logistic
  - ▶ For inner layers one may want symmetry/infinite range



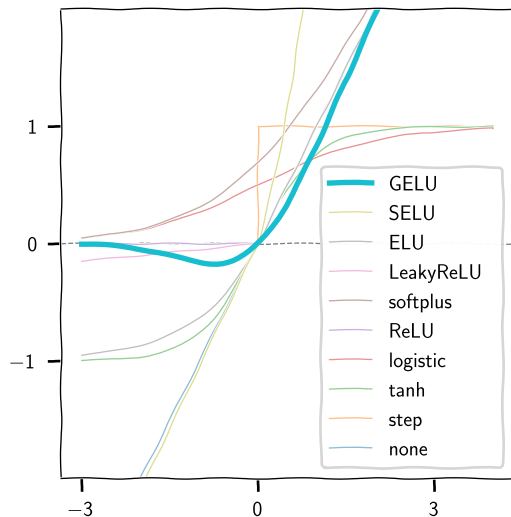
# Activation functions

- ▶ There is now a veritable zoo of activation functions
- ▶ Important properties to consider:
  - ▶ Smoothness
  - ▶ Saturation (at either end)
- ▶ Different roles depending on layer
  - ▶ At output it is useful to e.g. impose positivity with softmax or  $[0,1]$  boundedness with logistic
  - ▶ For inner layers one may want symmetry/infinite range



# Activation functions

- ▶ There is now a veritable zoo of activation functions
- ▶ Important properties to consider:
  - ▶ Smoothness
  - ▶ Saturation (at either end)
- ▶ Different roles depending on layer
  - ▶ At output it is useful to e.g. impose positivity with softmax or  $[0,1]$  boundedness with logistic
  - ▶ For inner layers one may want symmetry/infinite range



# Typical architectures

All problem dependent, but some reasonable guidelines are

## Regression

# input neurons	1 per input feature
# hidden layers	1 to 5
# neurons per hidden layer	Typically 10 to 100
# output neurons	1 per output dimension
Hidden activation	ReLU (or SELU)
Output activation	None
(positive outputs)	ReLU/Softplus
(bounded outputs)	Logistic/Tanh
Loss function	MSE
(if outliers)	MAE/Huber

## Classification

Same as Regression, except

- ▶ Loss function: cross-entropy
- ▶ # output neurons: same as number of labels/classes
- ▶ Output layer activation: Logistic for binary classifications and softmax for multiclass.

## Number of hidden layers

- ▶ Often a single layer will do (UAP)
- ▶ Deep networks allow you to do feature learning in the earlier layers
- ▶ This also enables transfer learning
- ▶ Start with one or two hidden layers, and gradually ramp up until you start overfitting.
- ▶ It may be helpful to use pre-trained networks

## Width of hidden layers

- ▶ Historically we “ramped down”, e.g. starting with 300, then 200, then 100
- ▶ In practice this makes little difference and adds tuning parameters
- ▶ Rectangular networks therefore more common
- ▶ Start with a small number and ramp up until the model starts overfitting

# Building neural networks

- ▶ Sequential API
  - ▶ Straightforward models
- ▶ Functional API
  - ▶ Complicated models
- ▶ Subclassing API
  - ▶ Dynamic models



## Example 3: Training Neural Networks

Go back to the notebook `training_ml_models.ipynb`

# Difficulties in training deep networks

- ▶ Exploding/vanishing gradients
- ▶ Not enough training
- ▶ Slow training
- ▶ Overfitting due to too many parameters

# Vanishing/Exploding gradients: weight initialisation

- ▶ Gradients vanish at plateaus, and explode if able to grow without bound (more common in recurrent neural networks)
- ▶ A step-change improvement in performance can be found by weight initialisation ( $\theta_0$ )
- ▶ Standard normally distributed (mean 0 variance 1) weights piped into an activation function do not result in mean 0 variance 1 outputs.
- ▶ A standard initialised network starts from a point of saturation.
- ▶ Terminology:  $\text{fan}_{\text{avg}} = \frac{1}{2}(\text{fan}_{\text{in}} + \text{fan}_{\text{out}})$ , where  $\text{fan}_{\text{in}} \equiv$  circuit terminology to describe the number of inputs to a layer

## Glorot

Activation: None, Tanh, logistic, softmax.

initialise weights: Normal with variance

$$\sigma^2 = \frac{1}{\text{fan}_{\text{avg}}} \text{ (or Uniform in } \pm\sqrt{3/\text{fan}_{\text{avg}}}\text{)}$$

## He

Activation: ReLU *et al.*

initialise weights:

$$\text{variance } \sigma^2 = \frac{2}{\text{fan}_{\text{in}}}$$

## Lecun

Activation: SELU

initialise weights:

$$\text{variance } \sigma^2 = \frac{1}{\text{fan}_{\text{in}}}$$

## Vanishing/Exploding gradients: Nonsaturating activation functions

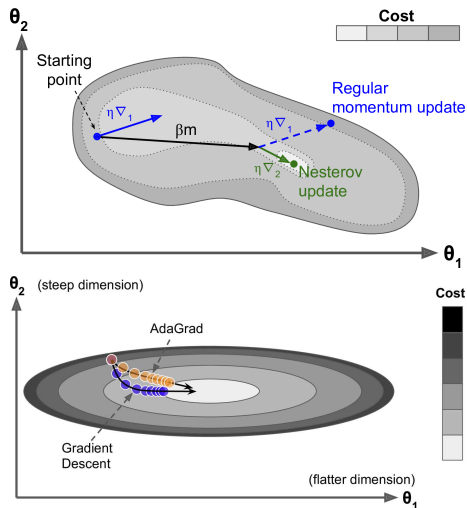
- ▶ Another way to fix the saturation problem is to choose a nonsaturating activation function (N.B. This is not how mother nature does it).
- ▶  $\text{ReLU} = \max(0, z)$  does not saturate, although it has a vanishing gradient by definition for negative inputs.
- ▶  $\text{LeakyReLU}_\alpha = \max(\alpha z, z)$  solves this.  $\alpha$  can be viewed as a hyperparameter (0.2, rand large leak or 0.01 for a small leak), randomised during training, or even fit for along the other parameters.
- ▶ ELU (exponential linear unit) is another choice
- ▶ Finally SELU are self-normalising ELUs. Very modern.
- ▶ GELU also trendy

## Vanishing/Exploding gradients: Batch normalisation

- ▶ Insert a normalisation step (zero centering and normalizing each input using a minibatch)
- ▶ At testing we use an exponential moving average over training for the shift parameters.
- ▶ Removes the need for standardisation if first layer is a BN layer
- ▶ Very much state-of-the-art
- ▶ Can be slower than ELU + He

# Faster optimisers

- ▶ Momentum optimization
  - ▶ Give some momentum/gradient memory to the trajectory
- ▶ Nesterov Accelerate Gradient
  - ▶ Use gradients ahead of the path to nudge
- ▶ AdaGrad
  - ▶ Scales gradients along steepest directions (measured by accumulated variance in components of gradient)
  - ▶ Often stops too early due to aggressive downscaling
- ▶ RMSProp
  - ▶ changes accumulation to an exponential moving average
- ▶ Adam & Nadam
  - ▶ Adaptive moment estimation
  - ▶ Combination of Momentum and RMSProp



# Learning rate scheduling

- ▶ Power scheduling:  $\eta(t) = \eta_0 / (1 + t/k)^c$
- ▶ Exponential scheduling:  $\eta(t) = \eta_0 0.1^{t/s}$
- ▶ Piecewise constant scheduling
- ▶ Performance scheduling: continuous version of early stopping

# Regularisation

## ▶ $\ell_1$ and $\ell_2$ regularisation

- ▶ These explicitly constrain the weights, and are easy to apply e.g.

```
layer = keras.layers.Dense(100, activation="elu",  
                             kernel_initializer="he_normal",  
                             kernel_regularizer=keras.regularizers.l2(0.01))
```

(try it out!  $\ell_1$  and  $\ell_{1,2}$  as well)

## ▶ Dropout

- ▶ At every training step, every neuron is dropped/zeroed in the calculation with probability  $p$
- ▶ No dropping after training (but minor correction of  $1 - p$  keep probability rescaling)
- ▶ Simple but shockingly effective

## ▶ MCDropout

- ▶ Cambridge-based theory paper linking dropout networks with ABC [arxiv:1506.02142]
- ▶ Puts monte-carlo dropping back into training to get errors

## ▶ Max-Norm regularisation

- ▶ Constrain the weights so  $\|w\|_2 < r$ , clipping if needed  $w \rightarrow wr/\|w\|_2$



# DeepNet guidelines (Géron)

Hyperparameter	Default value
Kernel initializer:	LeCun initialization
Activation function:	SELU
Normalization:	None (self-normalization)
Regularization:	Early stopping
Optimizer:	Nadam
Learning rate schedule:	Performance scheduling

Don't forget to standardize the input features!




- ▶ If you need sparsity try  $\ell_1 \pm$  FTRL optimisation + BN
- ▶ If you need low-latency, use fewer layers, avoid BN, SELU  $\rightarrow$  ReLU, consider sparsity & reducing precision
- ▶ If you are risk-sensitive consider MCDropout for performance boost and uncertainty

- ▶ If self-normalising & overfitting  
add  $\alpha$ -dropout  
(do not use other regularisation)
- ▶ If cannot self normalise  
try ELU instead of SELU  
(change initialisation)  
use BN after every hidden layer  
try max-norm or  $\ell_2$  regularisation

Happy training!

# Summary

- ▶ The data scientist's Python stack  
numpy, scipy, pandas, matplotlib, sklearn, +Keras/TensorFlow/PyTorch
- ▶ Principles & challenges of machine learning
- ▶ Theory of training, validation & testing
- ▶ Gradient descent and its pitfalls
- ▶ Regularisation and early stopping
- ▶ Neural networks Anatomy & training
- ▶ Recent advances in deep learning  
initialisation, activation, normalisation, optimisation, regularisation



## What we didn't cover!

- ▶ Choosing features
- ▶ Visualisation
- ▶ Classification